

EDA_on_Vehical_Insurance

October 22, 2024

Let's start with importing the libraries

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
```

This step is to import the data.

```
[2]: Data = pd.read_csv('/content/Vehicle_Insurance.csv')
```

```
[3]: # This is to view the first few rows of the dataset
Data.head()
```

```
[3]:
```

	id	Gender	Age	Driving_License	Region_Code	Previously_Insured	\
0	1	Male	44	1	28.0	0	
1	2	Male	76	1	3.0	0	
2	3	Male	47	1	28.0	0	
3	4	Male	21	1	11.0	1	
4	5	Female	29	1	41.0	1	

	Vehicle_Age	Vehicle_Damage	Annual_Premium	Policy_Sales_Channel	Vintage	\
0	> 2 Years	Yes	40454.0	26.0	217.0	
1	1-2 Year	No	33536.0	26.0	183.0	
2	> 2 Years	Yes	38294.0	26.0	27.0	
3	< 1 Year	No	28619.0	152.0	203.0	
4	< 1 Year	No	27496.0	152.0	39.0	

	Response
0	1.0
1	0.0
2	1.0
3	0.0
4	0.0

This step is to find the sum of null values in the Data set

```
[4]: Data.isnull().sum()
```

```
[4]: id          0
      Gender      0
      Age         0
      Driving_License  0
      Region_Code  0
      Previously_Insured  0
      Vehicle_Age    1
      Vehicle_Damage  1
      Annual_Premium  1
      Policy_Sales_Channel  1
      Vintage        1
      Response       1
      dtype: int64
```

This step is to get the mean, median and etc. for the data set.

```
[5]: Data.describe()
```

```
[5]:
```

	id	Age	Driving_License	Region_Code \
count	354405.000000	354405.000000	354405.000000	354405.000000
mean	177203.000000	38.813603	0.997884	26.397531
std	102308.055414	15.500885	0.045954	13.232660
min	1.000000	20.000000	0.000000	0.000000
25%	88602.000000	25.000000	1.000000	15.000000
50%	177203.000000	36.000000	1.000000	28.000000
75%	265804.000000	49.000000	1.000000	35.000000
max	354405.000000	85.000000	1.000000	52.000000

	Previously_Insured	Annual_Premium	Policy_Sales_Channel \
count	354405.000000	354404.000000	354404.000000
mean	0.457937	30556.903475	112.085408
std	0.498228	17224.081990	54.174460
min	0.000000	2630.000000	1.000000
25%	0.000000	24395.000000	29.000000
50%	0.000000	31660.000000	134.000000
75%	1.000000	39391.000000	152.000000
max	1.000000	540165.000000	163.000000

	Vintage	Response
count	354404.000000	354404.000000
mean	154.362620	0.122961
std	83.664171	0.328393
min	10.000000	0.000000
25%	82.000000	0.000000
50%	154.000000	0.000000
75%	227.000000	0.000000
max	299.000000	1.000000

This step gives us information regarding the whole data set.

```
[6]: Data.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 354405 entries, 0 to 354404
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    354405 non-null  int64
1   Gender                354405 non-null  object
2   Age                   354405 non-null  int64
3   Driving_License       354405 non-null  int64
4   Region_Code           354405 non-null  float64
5   Previously_Insured    354405 non-null  int64
6   Vehicle_Age           354404 non-null  object
7   Vehicle_Damage        354404 non-null  object
8   Annual_Premium        354404 non-null  float64
9   Policy_Sales_Channel  354404 non-null  float64
10  Vintage                354404 non-null  float64
11  Response               354404 non-null  float64
dtypes: float64(5), int64(4), object(3)
memory usage: 32.4+ MB
```



```
[7]: # To remove duplicate rows based on all columns
Data.drop_duplicates(inplace=True)
```



```
[8]: Data.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 354405 entries, 0 to 354404
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    354405 non-null  int64
1   Gender                354405 non-null  object
2   Age                   354405 non-null  int64
3   Driving_License       354405 non-null  int64
4   Region_Code           354405 non-null  float64
5   Previously_Insured    354405 non-null  int64
6   Vehicle_Age           354404 non-null  object
7   Vehicle_Damage        354404 non-null  object
8   Annual_Premium        354404 non-null  float64
9   Policy_Sales_Channel  354404 non-null  float64
10  Vintage                354404 non-null  float64
11  Response               354404 non-null  float64
dtypes: float64(5), int64(4), object(3)
memory usage: 32.4+ MB
```



```
[9]: # Drawing boxplot to identify the outliers in the data set
for column in Data.columns:
    if pd.api.types.is_numeric_dtype(Data[column]):
        plt.figure()
        sns.boxplot(Data[column])
```





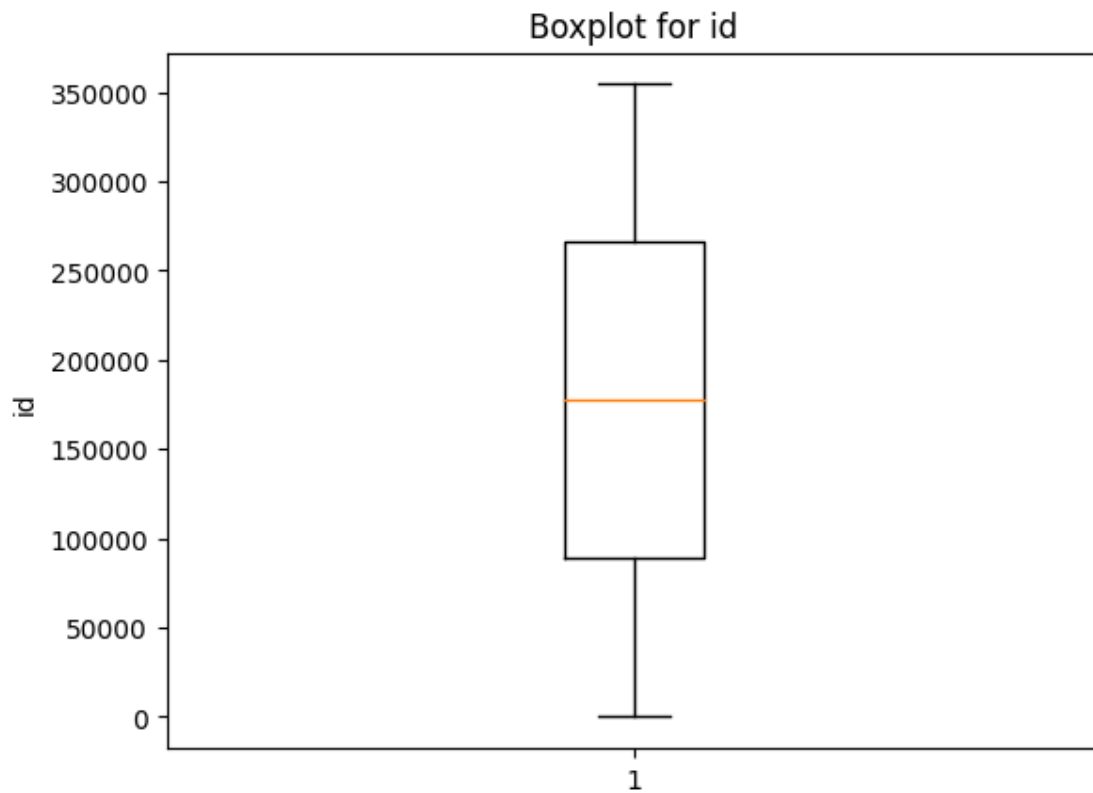





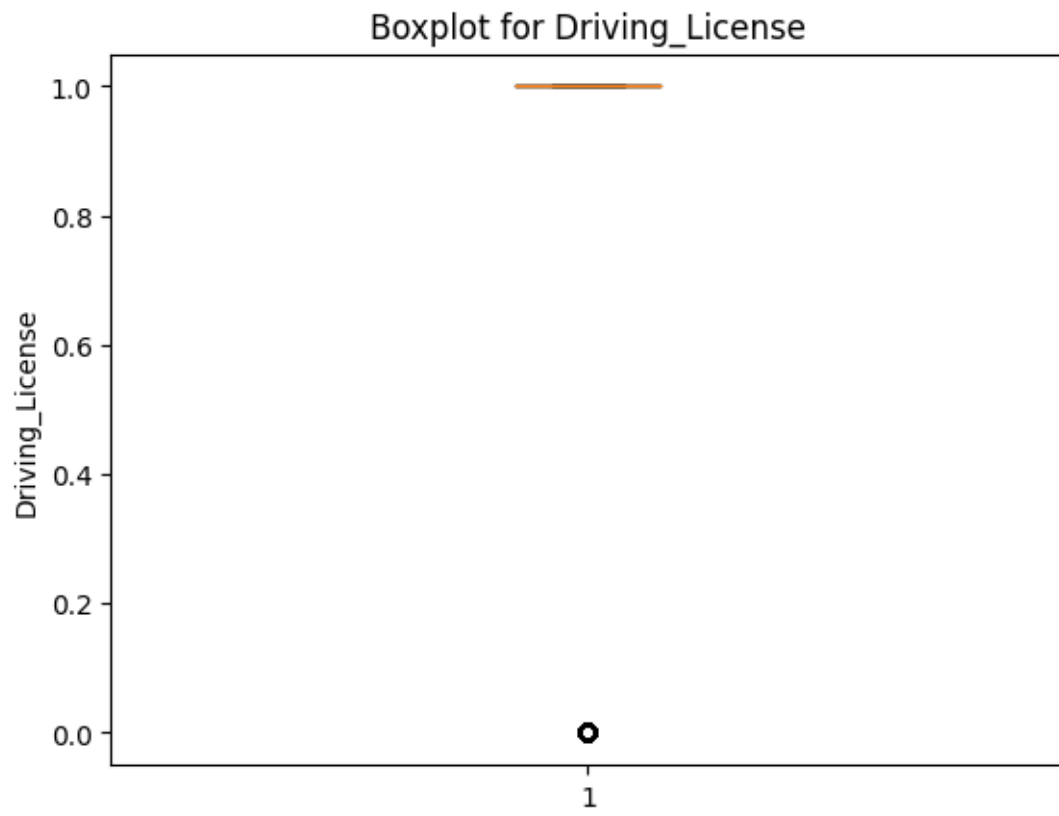




```
[11]: # Drawing the boxplots to identify the outliers in all columns.
```

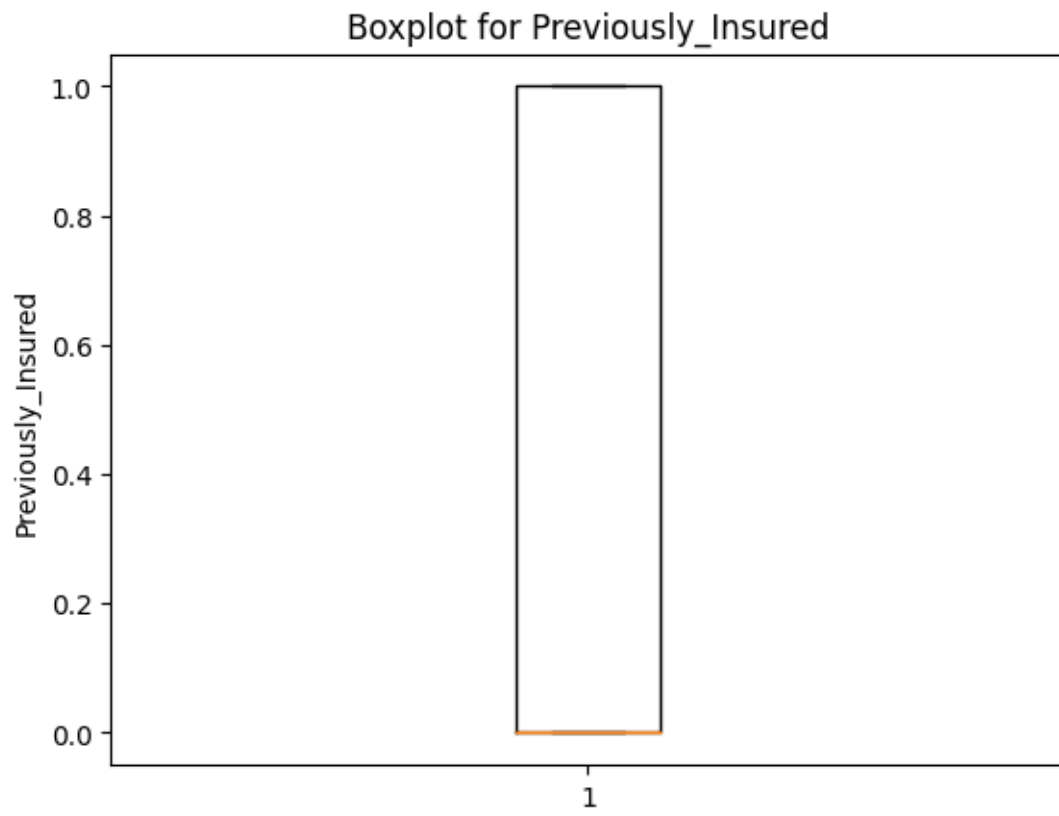
```
for column in Data.columns:  
    if pd.api.types.is_numeric_dtype(Data[column]):  
        plt.figure()  
        plt.boxplot(Data[column])  
        plt.title(f"Boxplot for {column}")  
        plt.ylabel(column)  
        plt.show()
```

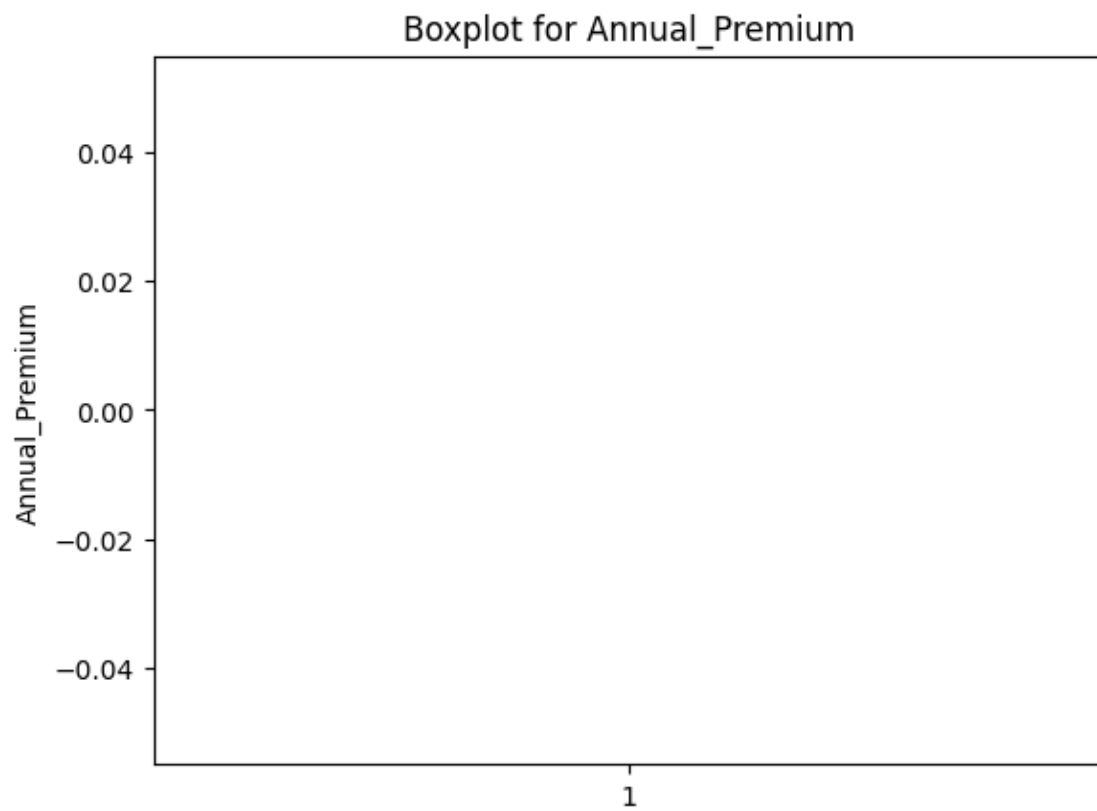


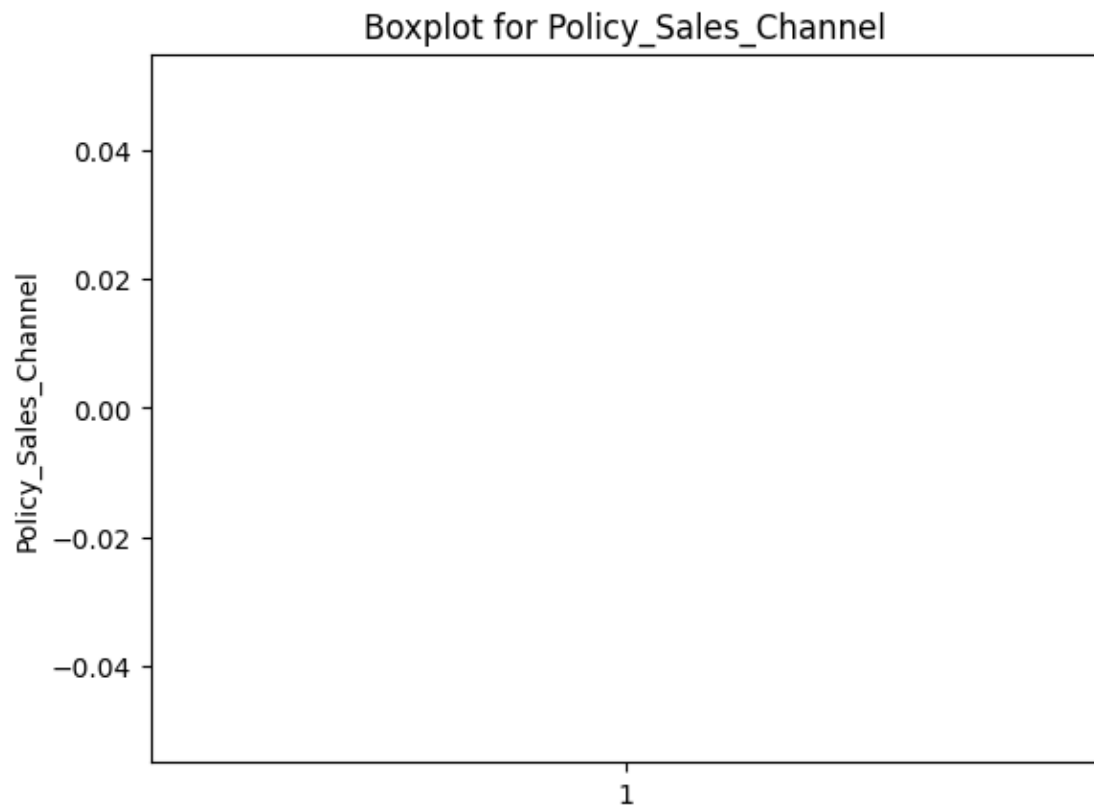


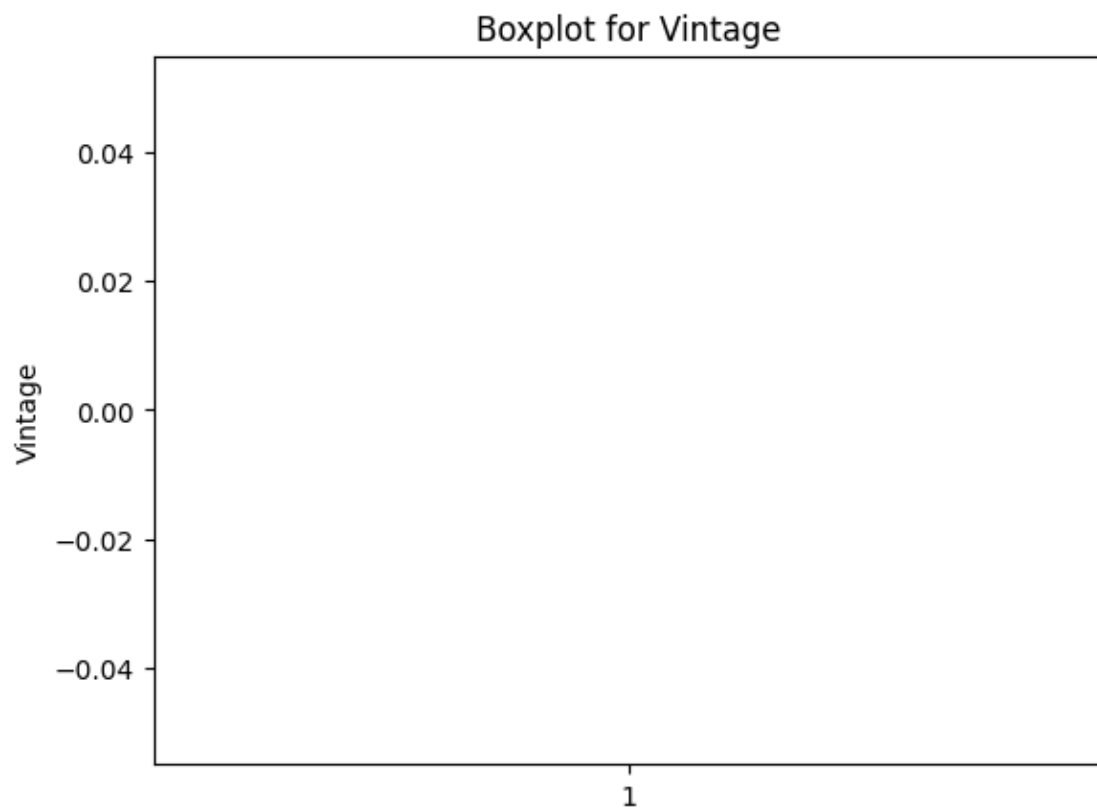


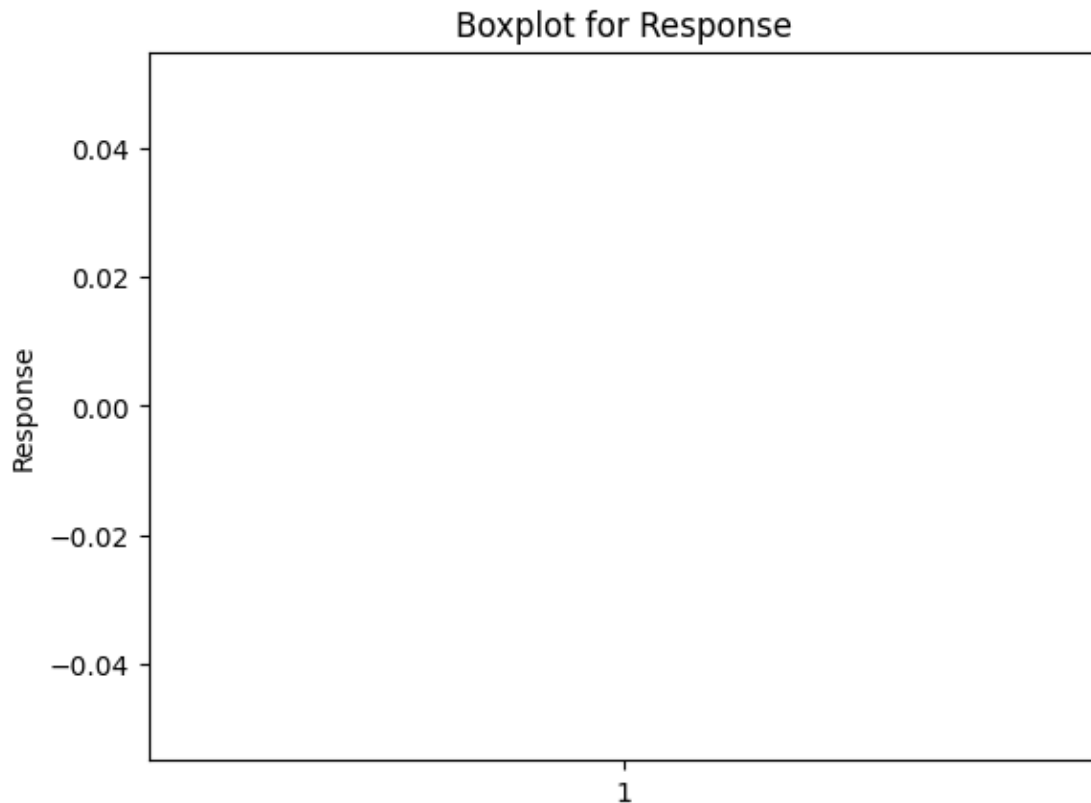












[12]: *#Replacing the outliers in the Annual Premium column with the median.*

```
def replace_outliers_with_median(df, column):
    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    df[column] = np.where((df[column] < lower_bound) | (df[column] >
↪upper_bound),
                        df[column].median(), df[column])


    return df


Data = replace_outliers_with_median(Data, 'Annual_Premium')
```

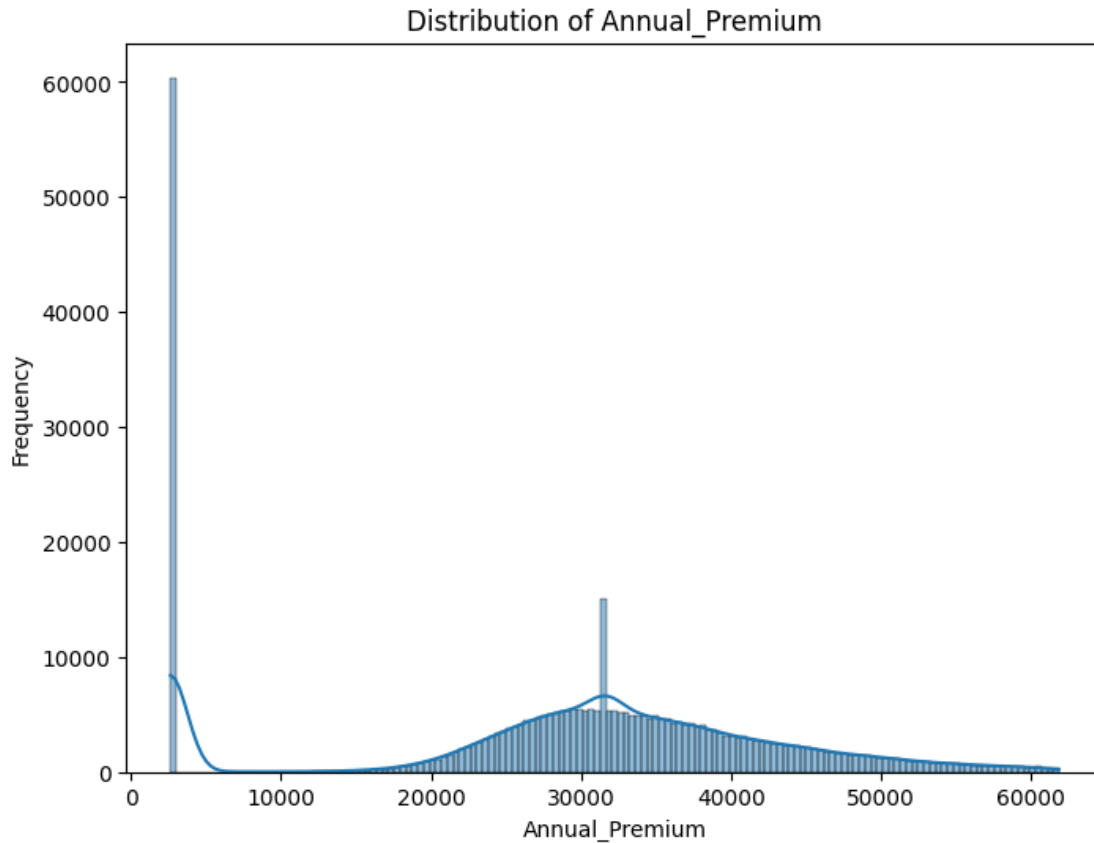
Since the outliers and the missing values are handled now we are moving forward with the Data Visualization.

[13]: *# Now we are utilizing various visualization techniques to explore the ↪distribution of key variables.*

```
# Histograms for numerical features
for col in ['Age', 'Vintage', 'Annual_Premium']:
    plt.figure(figsize=(8, 6))
    sns.histplot(Data[col], kde=True)
    plt.title(f'Distribution of {col}')
    plt.xlabel(col)
    plt.ylabel('Frequency')
    plt.show()
```






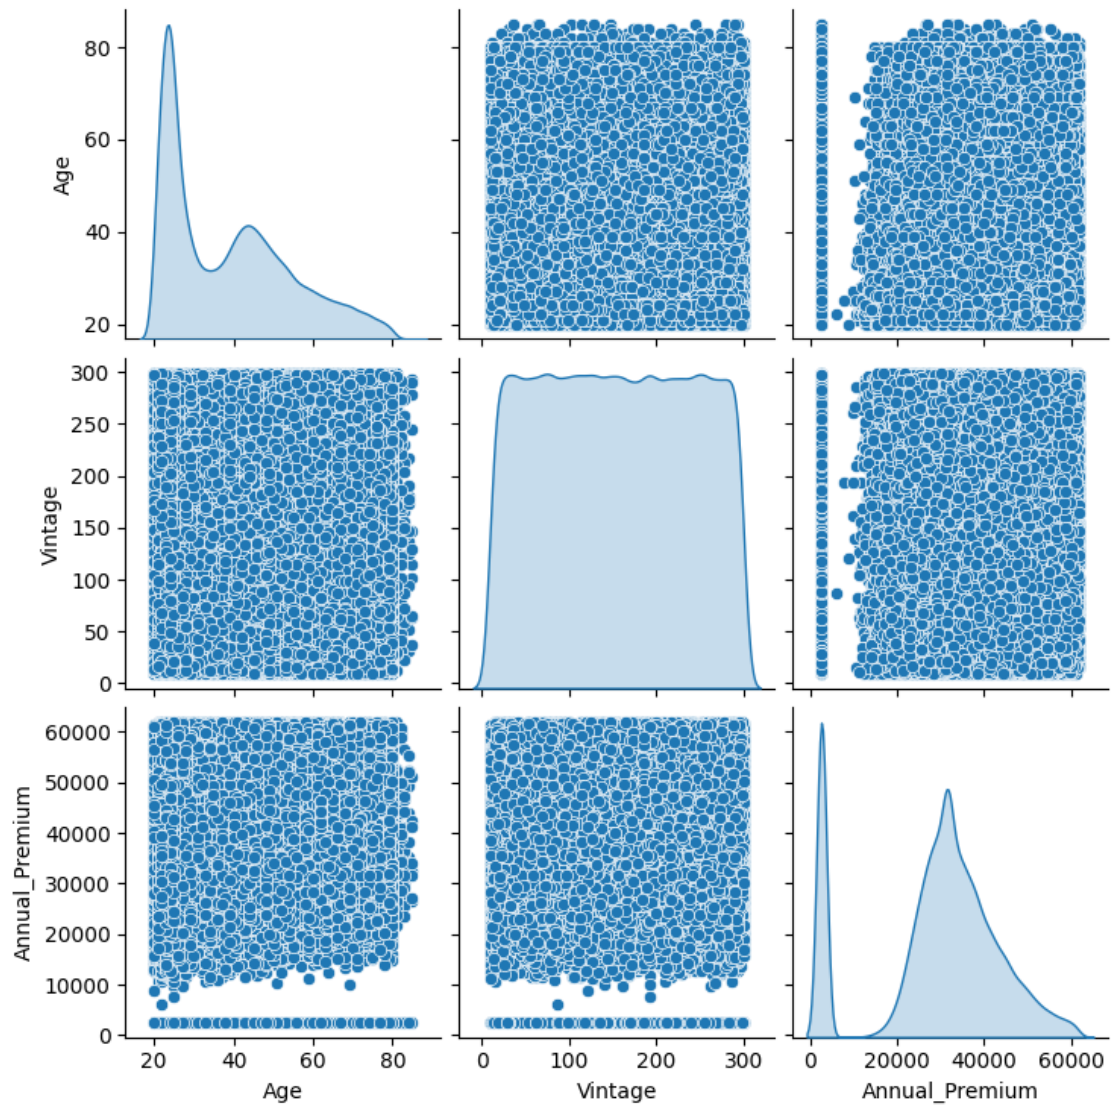


```
[14]: # Scatter plots to explore relationships between variables
plt.figure(figsize=(10, 6))
sns.scatterplot(x='Age', y='Annual_Premium', data=Data, hue='Vehicle_Damage')
plt.title('Age vs. Annual Premium')
plt.show()
```


```
/usr/local/lib/python3.10/dist-packages/IPython/core/pylabtools.py:151:
UserWarning: Creating legend with loc="best" can be slow with large amounts of
data.
  fig.canvas.print_figure(bytes_io, **kw)
```

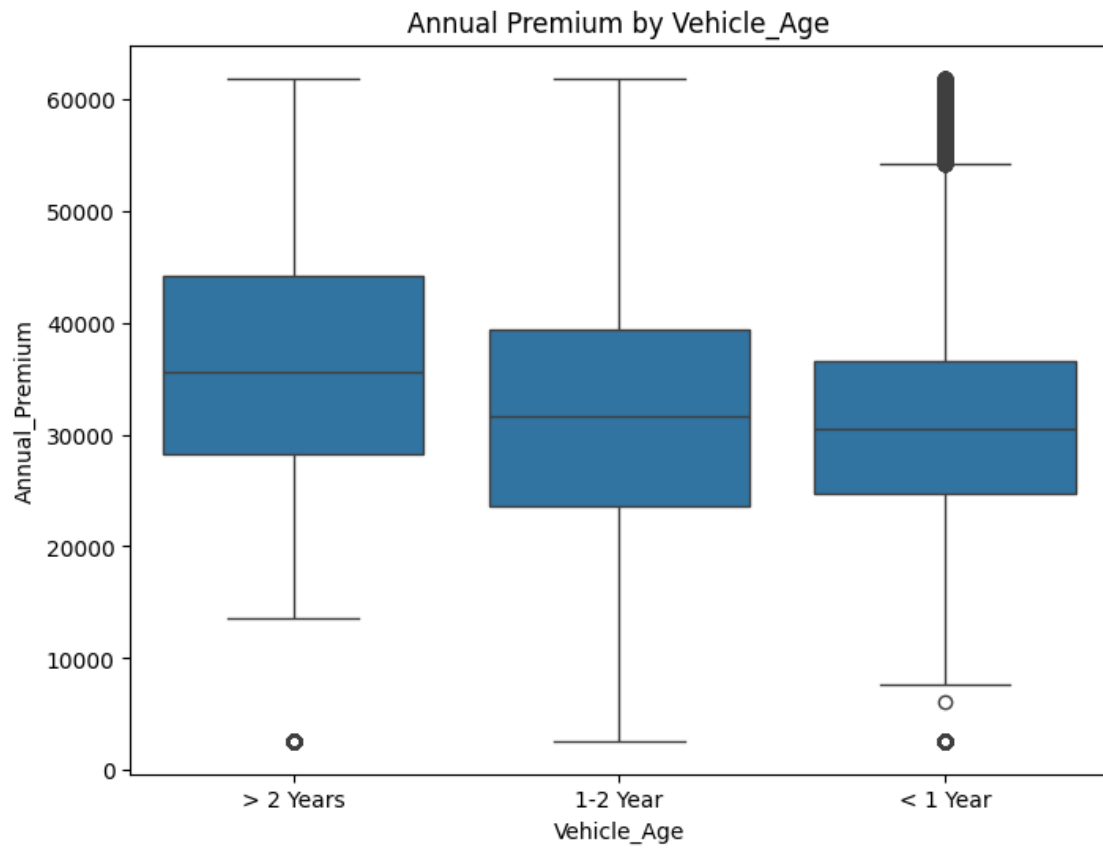


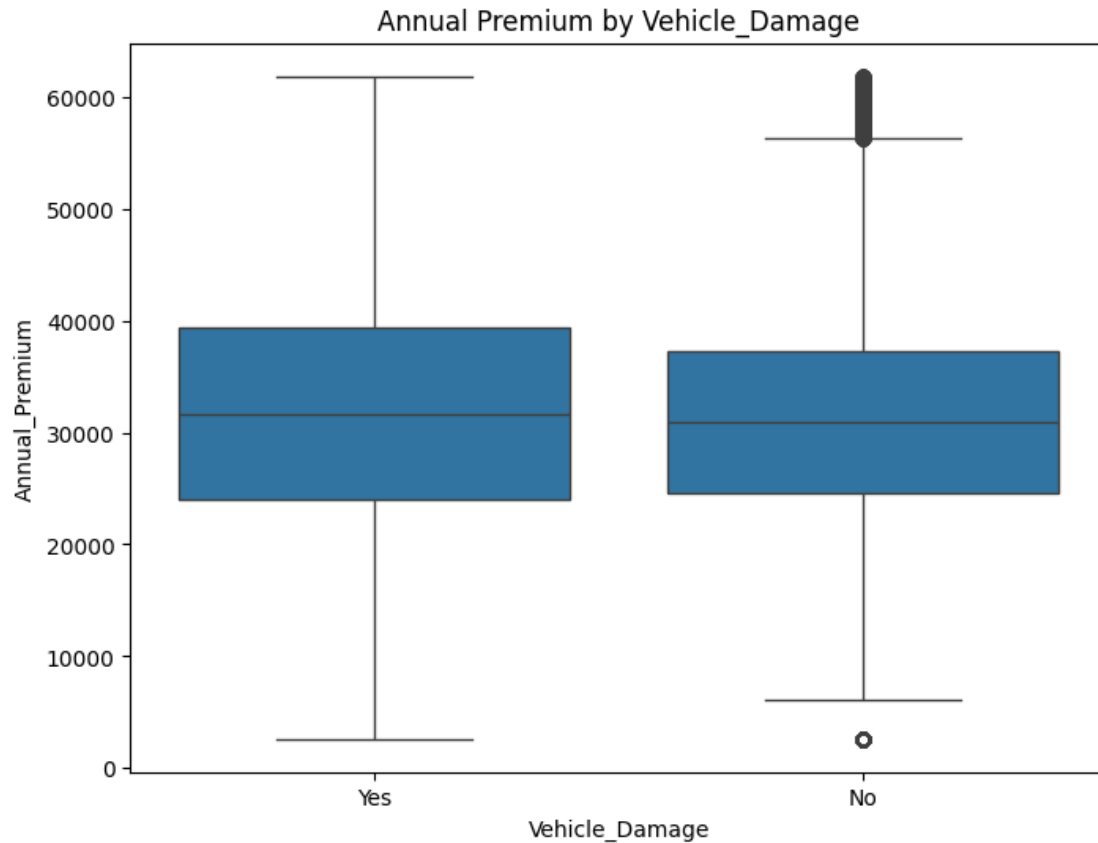
```
[15]: # Pairplot for multiple numerical variables
sns.pairplot(Data[['Age', 'Vintage', 'Annual_Premium']], diag_kind='kde')
plt.show()
```




```
[16]: # Box plots for categorical features
for col in ['Gender', 'Vehicle_Age', 'Vehicle_Damage']:
    plt.figure(figsize=(8, 6))
    sns.boxplot(x=col, y='Annual_Premium', data=Data)
    plt.title(f'Annual Premium by {col}')
    plt.show()
```

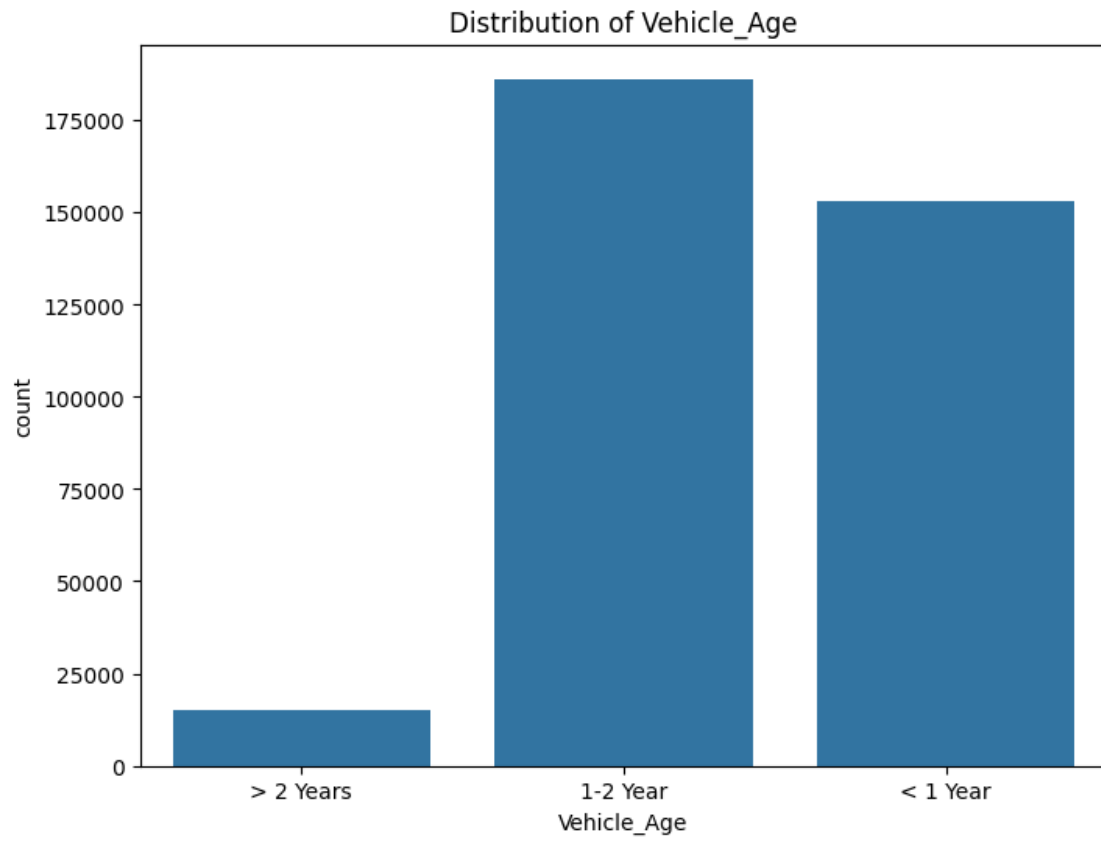


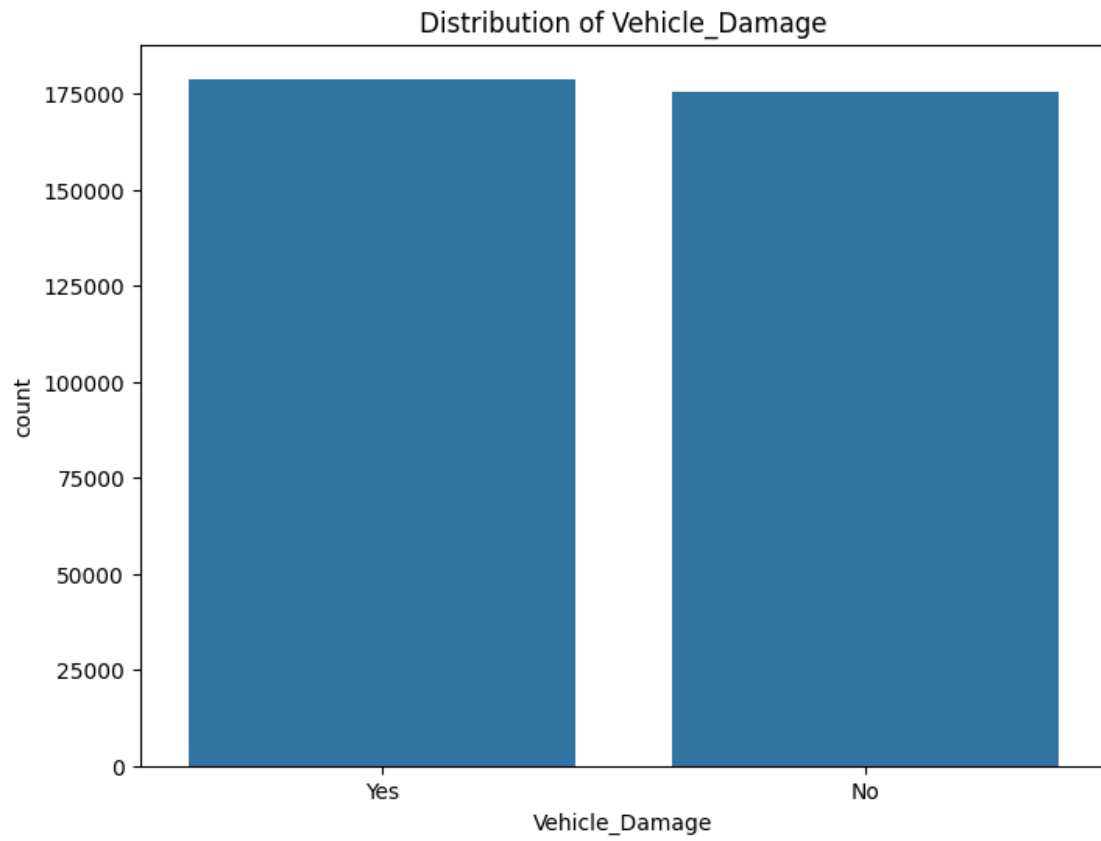


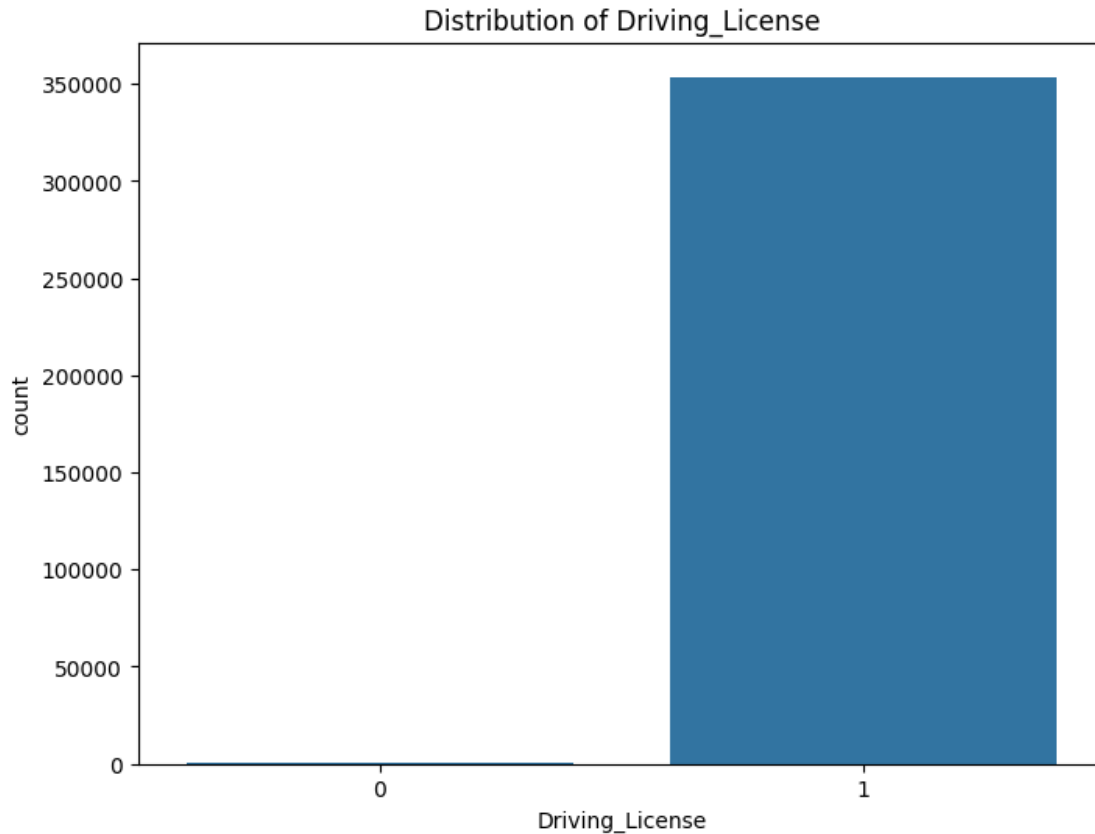


```
[17]: # Count plots for categorical features
for col in ['Gender', 'Vehicle_Age', 'Vehicle_Damage', 'Driving_License']:
    plt.figure(figsize=(8, 6))
    sns.countplot(x=col, data=Data)
    plt.title(f'Distribution of {col}')
    plt.show()
```









```
[18]: # Interactive plot using Plotly
fig = px.histogram(Data, x='Annual_Premium', color='Vehicle_Damage',
    ↪marginal='box',
                    title='Annual Premium Distribution by Vehicle Damage',
                    labels={'Annual_Premium': 'Annual Premium'})
fig.show()
```

Here's a summary of key insights from the data visualization code's:

1. **Data Cleaning:** The code begins by handling missing values and outliers. Outliers in the 'Annual_Premium' column are replaced with the median value. This is crucial for reliable analysis as outliers can skew results.
2. **Distribution of Numerical Variables:** Histograms reveal the distribution of 'Age', 'Vintage', and 'Annual_Premium'. These visualizations help understand the central tendency, spread, and skewness of these variables. The code also checks for normality using kernel density estimates (KDE).
3. **Relationship between Age and Annual Premium:** The scatter plot of 'Age' vs. 'Annual_Premium', colored by 'Vehicle_Damage', shows a potential correlation between age and annual premium and how vehicle damage status influences this relationship.
4. **Relationships between Numerical Variables:** The pairplot visualizes the relationships

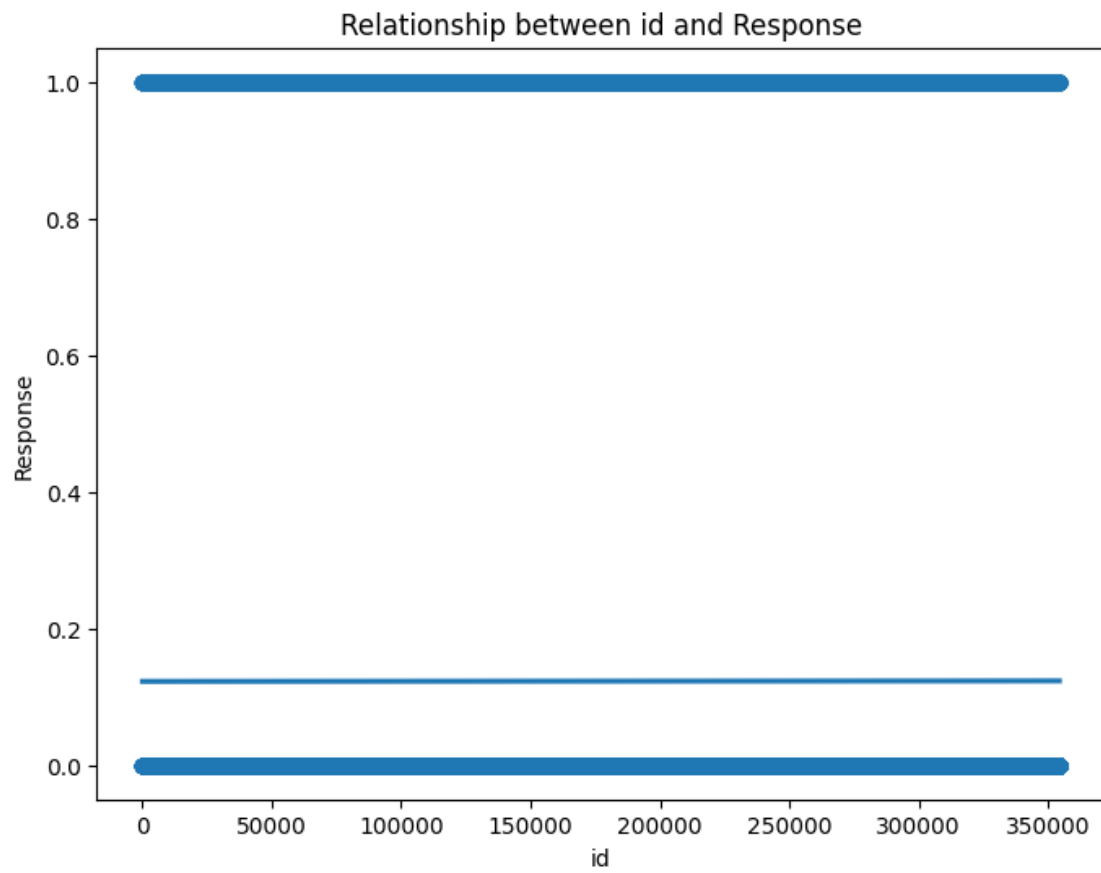
between 'Age', 'Vintage', and 'Annual_Premium', including their individual distributions (using KDE plots). This helps to identify potential correlations or dependencies between these variables.

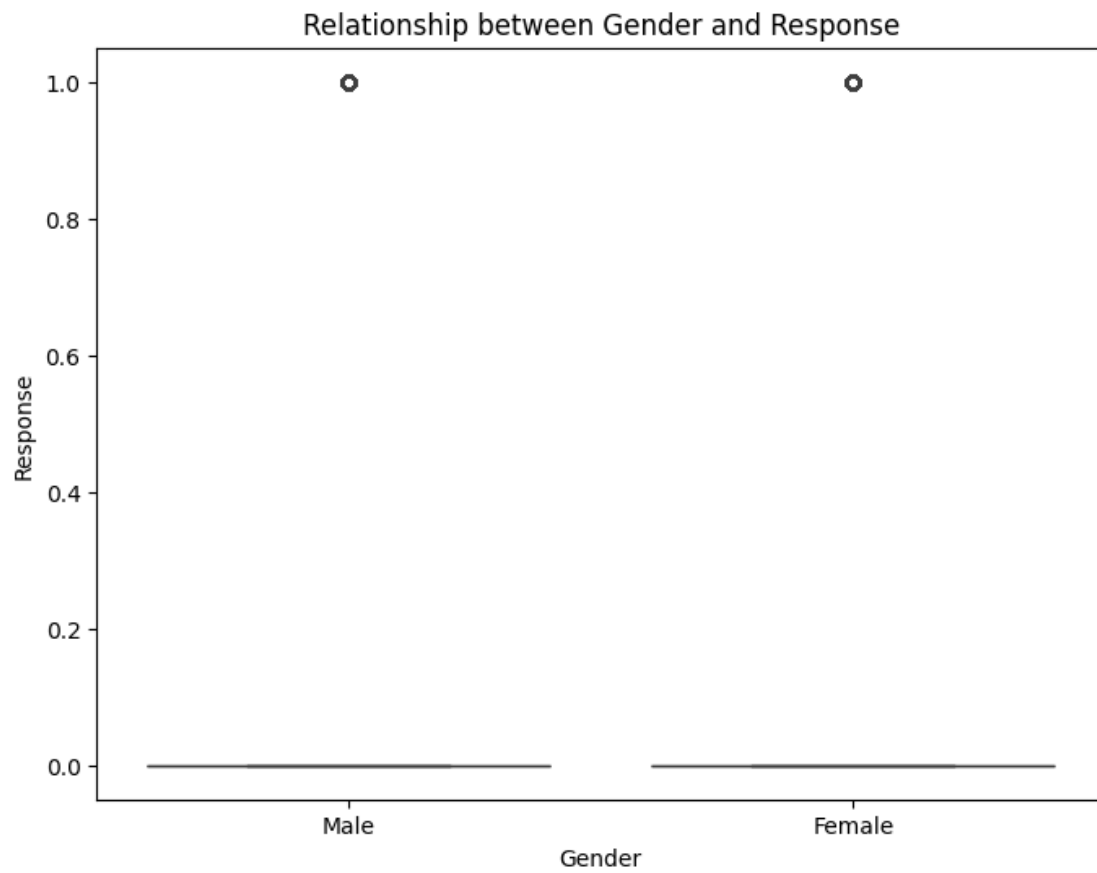
5. **Impact of Categorical Variables on Annual Premium:** Box plots illustrate how 'Gender', 'Vehicle_Age', and 'Vehicle_Damage' affect 'Annual_Premium'. They show the distribution of premiums within each category and help identify potential differences in premiums based on these factors.
6. **Distribution of Categorical Variables:** Count plots visualize the frequency of each category in 'Gender', 'Vehicle_Age', 'Vehicle_Damage', and 'Driving_License', giving an overview of the distribution of these categorical variables within the dataset.
7. **Interactive Visualization:** The Plotly histogram provides an interactive way to explore the distribution of 'Annual_Premium', broken down by 'Vehicle_Damage'. The inclusion of a box plot in the margin enhances the understanding of the distribution's central tendency, spread, and potential outliers.


Now the data vizualization step is completed. Let's move on to the Feature Analysis (Target variable - Insurace Claims or Response column)


```
[22]: # Analyze the relationship between features and the target variable ('Response')


for column in Data.columns:
    if column != 'Response': # Exclude the target variable itself
        if pd.api.types.is_numeric_dtype(Data[column]):
            plt.figure(figsize=(8, 6))
            sns.regplot(x=column, y='Response', data=Data)
            plt.title(f'Relationship between {column} and Response')
            plt.show()
        else: # Categorical features
            plt.figure(figsize=(8, 6))
            sns.boxplot(x=column, y='Response', data=Data)
            plt.title(f'Relationship between {column} and Response')
            plt.show()
```






























```
[25]: # Calculate the correlation between features and the target variable
# Select only numeric columns before calculating correlation
numeric_data = Data.select_dtypes(include=np.number)
correlation_with_response = numeric_data.corr()['Response'].drop('Response') #_
    ↳ Drop self correlation

print("Correlation with Response Variable:\n", correlation_with_response)
```

Correlation with Response Variable:

id	0.000654
Age	0.111417
Driving_License	0.010512
Region_Code	0.010702
Previously_Insured	-0.341637
Annual_Premium	0.019536
Policy_Sales_Channel	-0.140131
Vintage	-0.000445
Name: Response, dtype: float64	

```
[26]: # Visualize correlations with a bar plot
plt.figure(figsize=(10, 6))
correlation_with_response.plot(kind='bar')
plt.title('Correlation with Response Variable')
plt.xlabel('Features')
plt.ylabel('Correlation Coefficient')
plt.xticks(rotation=45, ha='right')
plt.show()
```



The feature analysis code investigates the relationship between various features in the dataset and the target variable, 'Response'. Here's a breakdown of the key insights:

1. **Visual Exploration of Feature Relationships:** The code iterates through each feature (excluding 'Response'). For numerical features, it uses regression plots (`sns.regplot`) to visualize the linear relationship with 'Response'. For categorical features, box plots (`sns.boxplot`) show the distribution of 'Response' for each category, revealing potential differences in the target variable based on the categorical feature.
2. **Correlation Analysis (Numerical Features):** The code calculates the correlation coefficients between numerical features and the 'Response' variable. This quantifies the linear association between each numerical feature and the target variable. Positive correlations suggest that as the feature value increases, the likelihood of a positive 'Response' also increases,

while negative correlations indicate the opposite. The strength of the correlation is indicated by the absolute value of the coefficient (closer to 1 indicates a stronger linear relationship). A bar plot then visualizes these correlations.

Interpreting the Results:

The visualizations and correlation coefficients provide crucial insights into which features might be important predictors of 'Response'. Strong correlations (positive or negative) suggest that the corresponding features are likely significant in predicting the outcome. Visualizations (especially the box plots for categorical variables) can also highlight non-linear relationships or interactions that might not be fully captured by correlation coefficients. For example, if a box plot for a categorical variable shows significantly different distributions of the target variable between categories, that variable is likely important even if the correlation coefficient is weak or non-existent.

1 Summarization

Here's a summary of the key processes and factors:

1. Data Loading and Initial Exploration:

- **Key Factors:** The process begins by loading the dataset and performing initial checks. `Data.head()`, `Data.isnull().sum()`, `Data.describe()`, and `Data.info()` provide crucial initial insights into the data's structure, summary statistics, and the presence of missing values.
- **Key Process:** The code imports necessary libraries (NumPy, Pandas, Matplotlib, Seaborn, Plotly) and reads the CSV file into a Pandas DataFrame.

2. Data Cleaning:

- **Key Factors:** Duplicate rows are removed, and outliers in the 'Annual_Premium' column are addressed by replacing them with the median. Outlier handling is vital to prevent skewed results.
- **Key Process:** `Data.drop_duplicates()` eliminates redundant entries. A function `replace_outliers_with_median` uses the IQR method to detect and replace outliers.

3. Data Visualization:

- **Key Factors:** Various visualization techniques are used. Histograms show variable distributions, scatter plots explore relationships between variables (e.g., age and annual premium), pair plots visualize multiple numerical variables' relationships. Box plots illustrate how categorical variables (gender, vehicle age, damage) influence the annual premium. Count plots show the frequency of categories. Plotly is used for interactive visualizations.
- **Key Process:** The code uses Seaborn and Matplotlib functions for static visualizations and Plotly for interactive ones. The goal is to understand the data's characteristics, patterns, and potential correlations.

4. Feature Analysis (Target Variable - 'Response'):

- **Key Factors:** The analysis focuses on the relationship between features and the target variable ('Response'), which likely represents insurance claims or a similar binary outcome. Regression plots for numerical features and box plots for categorical features visualize these

relationships. Correlation coefficients quantify the linear relationships between numerical features and 'Response', aiding in identifying potentially important predictors.

- **Key Process:** Correlation analysis and visualization help determine which features significantly impact the target variable.

Overall Highlights:

- **Comprehensive EDA:** The code demonstrates a systematic approach to EDA, covering data cleaning, exploration, and feature analysis.
- **Variety of Visualization Techniques:** The use of multiple visualization types provides a comprehensive understanding of the dataset.
- **Focus on Target Variable:** The code explicitly investigates how different features relate to the target variable, essential for predictive modeling.
- **Handling Missing Values and Outliers:** The code correctly addresses data quality issues, critical for meaningful analysis.
- **Interactive Visualization:** The use of Plotly adds an interactive dimension to the exploration, offering more flexibility and depth to the analysis.

This thorough analysis equips a data scientist to select relevant features and choose appropriate modeling techniques for the vehicle insurance dataset.